

The aim here is to see whether having a richer Wikipedia/Wikidata profile makes a brand more likely to get mentioned by ChatGPT when asked different neutral prompts.

Dataset Contains:

Prompts: 303

Brands: 2288

1) Importing Libraries

```
In [23]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler
from scipy.stats import spearmanr
import statmodels.api as sm
import statmodels.formula.api as smf
import seaborn as sns

In [2]: df = pd.read_csv("merged_output.csv", encoding='utf-8')
```

3) Quick data

```
In [4]: print("Rows:", len(df), " | Prompts:", df['Prompt'].nunique(), " | Brands (rows):", df['brand_Name'].nunique())
df[['PageExists']].value_counts()

Out [4]: Rows: 2763 | Prompts: 303 | Brands (rows): 2288
PageExists
True      1815
False     948
dtype: int64

In [3]: df.columns

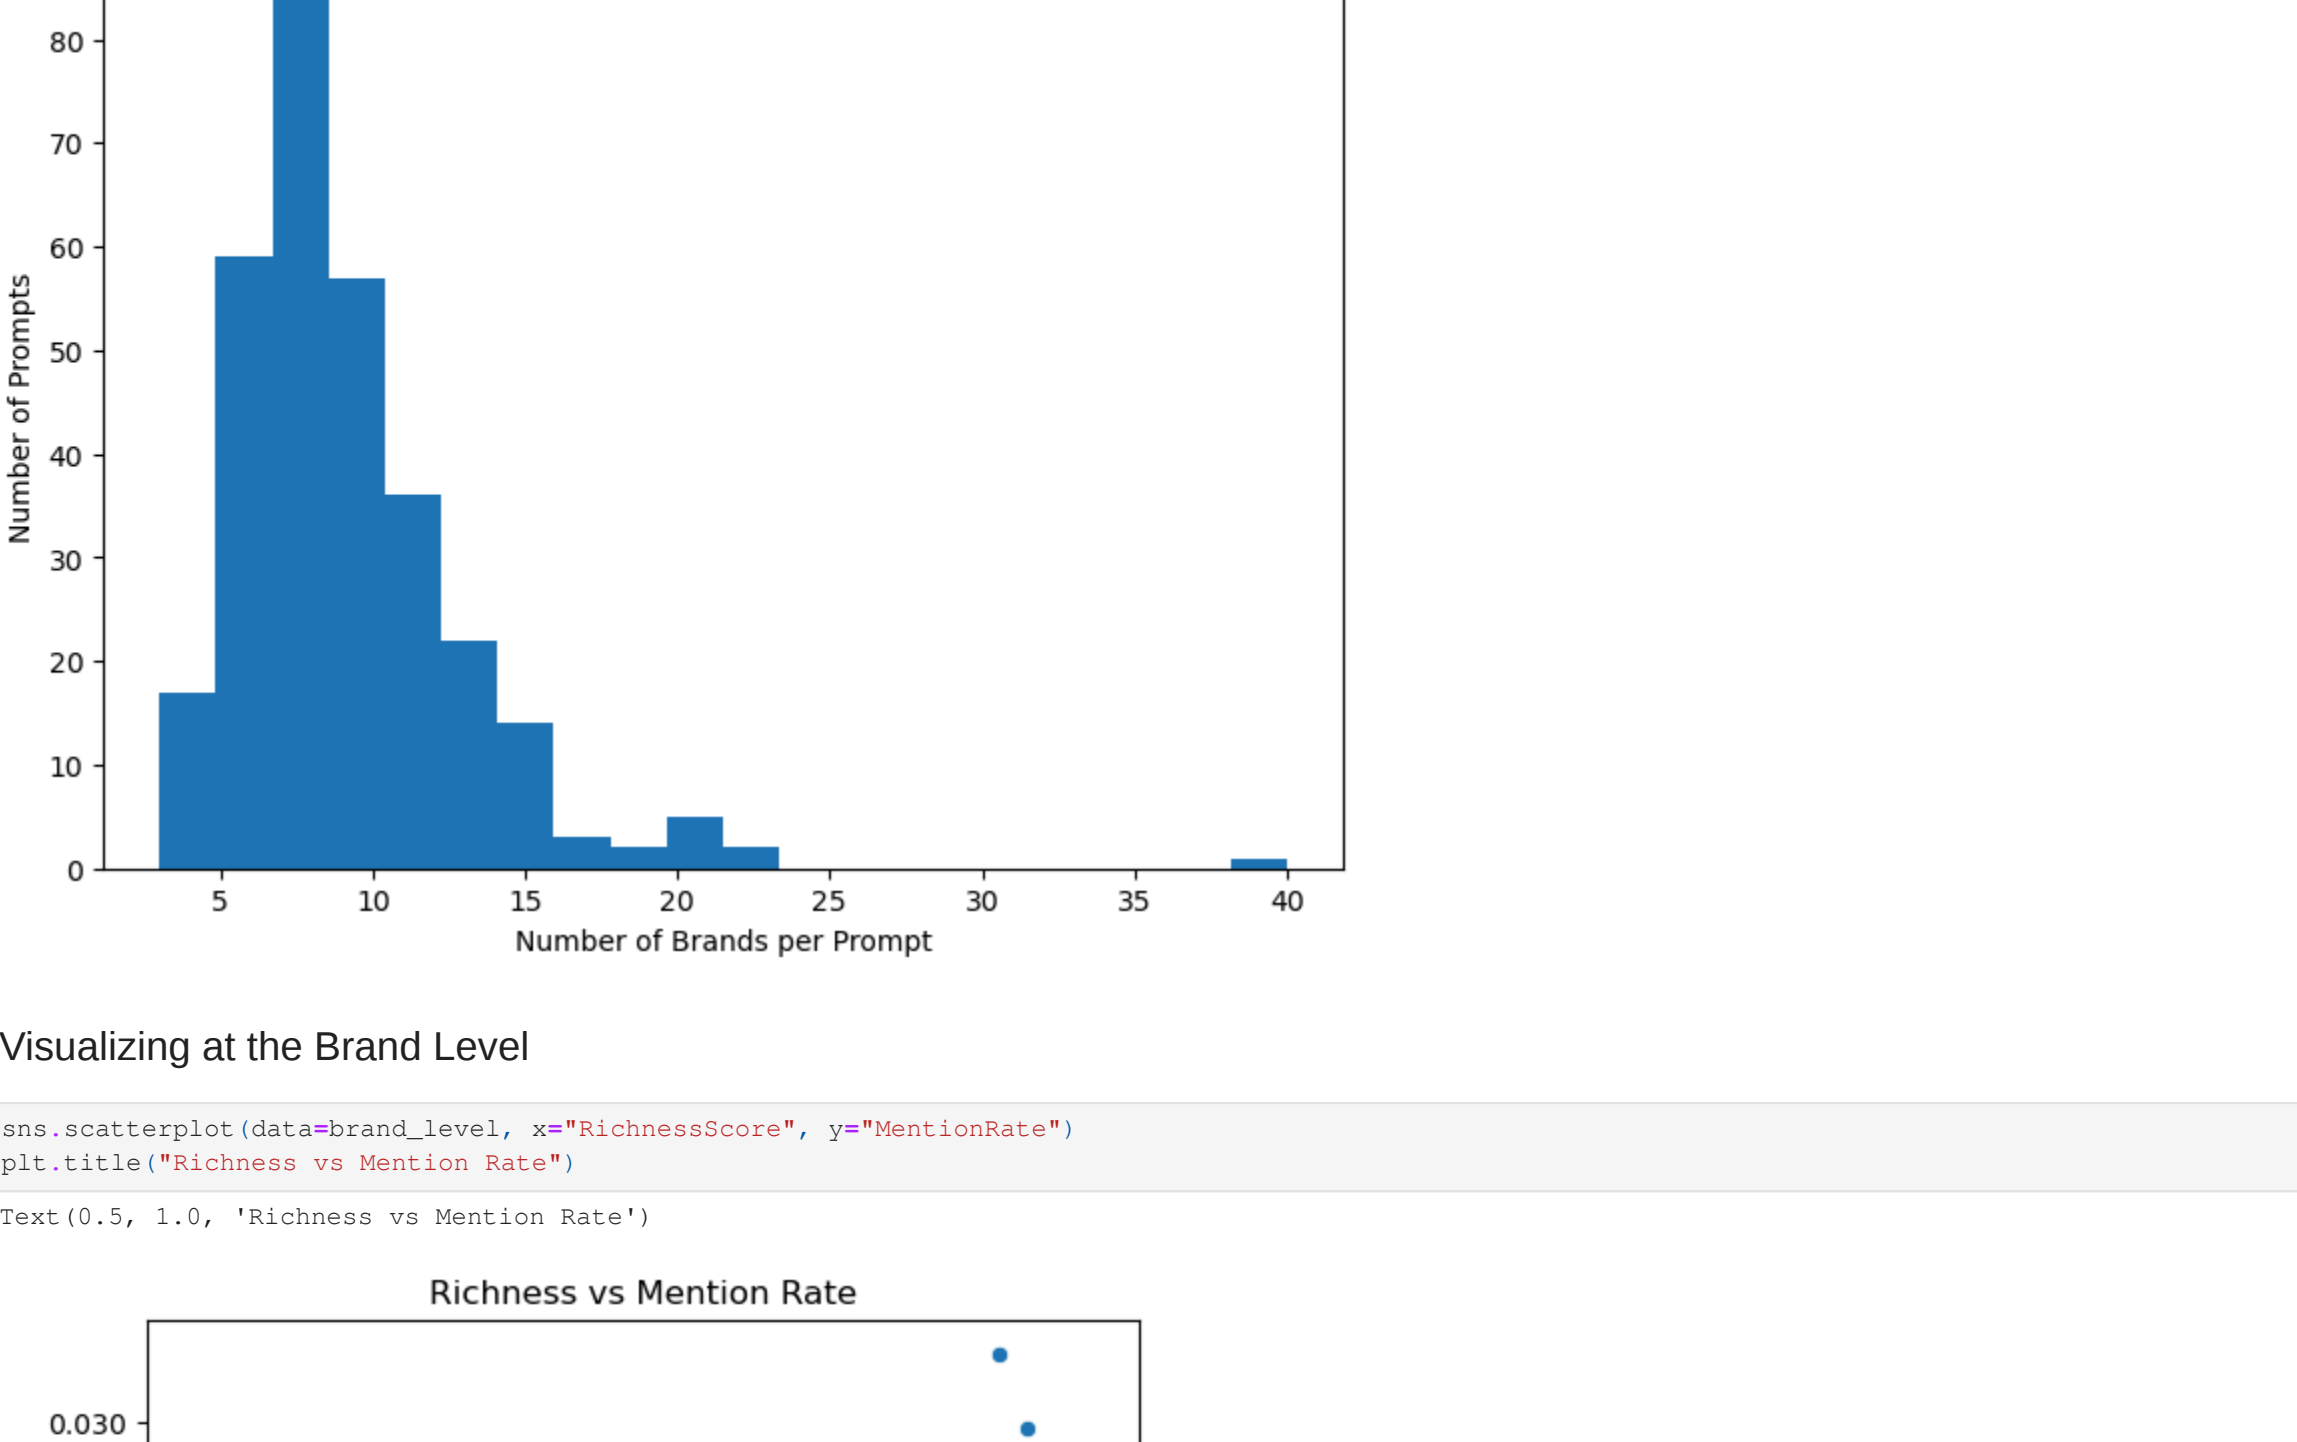
Out [5]: Index(['Prompt', 'brand_Name', 'mention_type', 'Rank', 'Brand', 'PageExists',
      'PageLength', 'RefCount', 'Categories', 'InternalLinks',
      'AvgPageviews90d', 'SiteLinks', 'Statements'],
      dtype='object')
```

Visualizing Data for Better Understanding


Top 10 Most Frequent Brands:




Mention Type Split (Ranked vs Bullet):




Distribution of Brands Mentioned per Prompt:



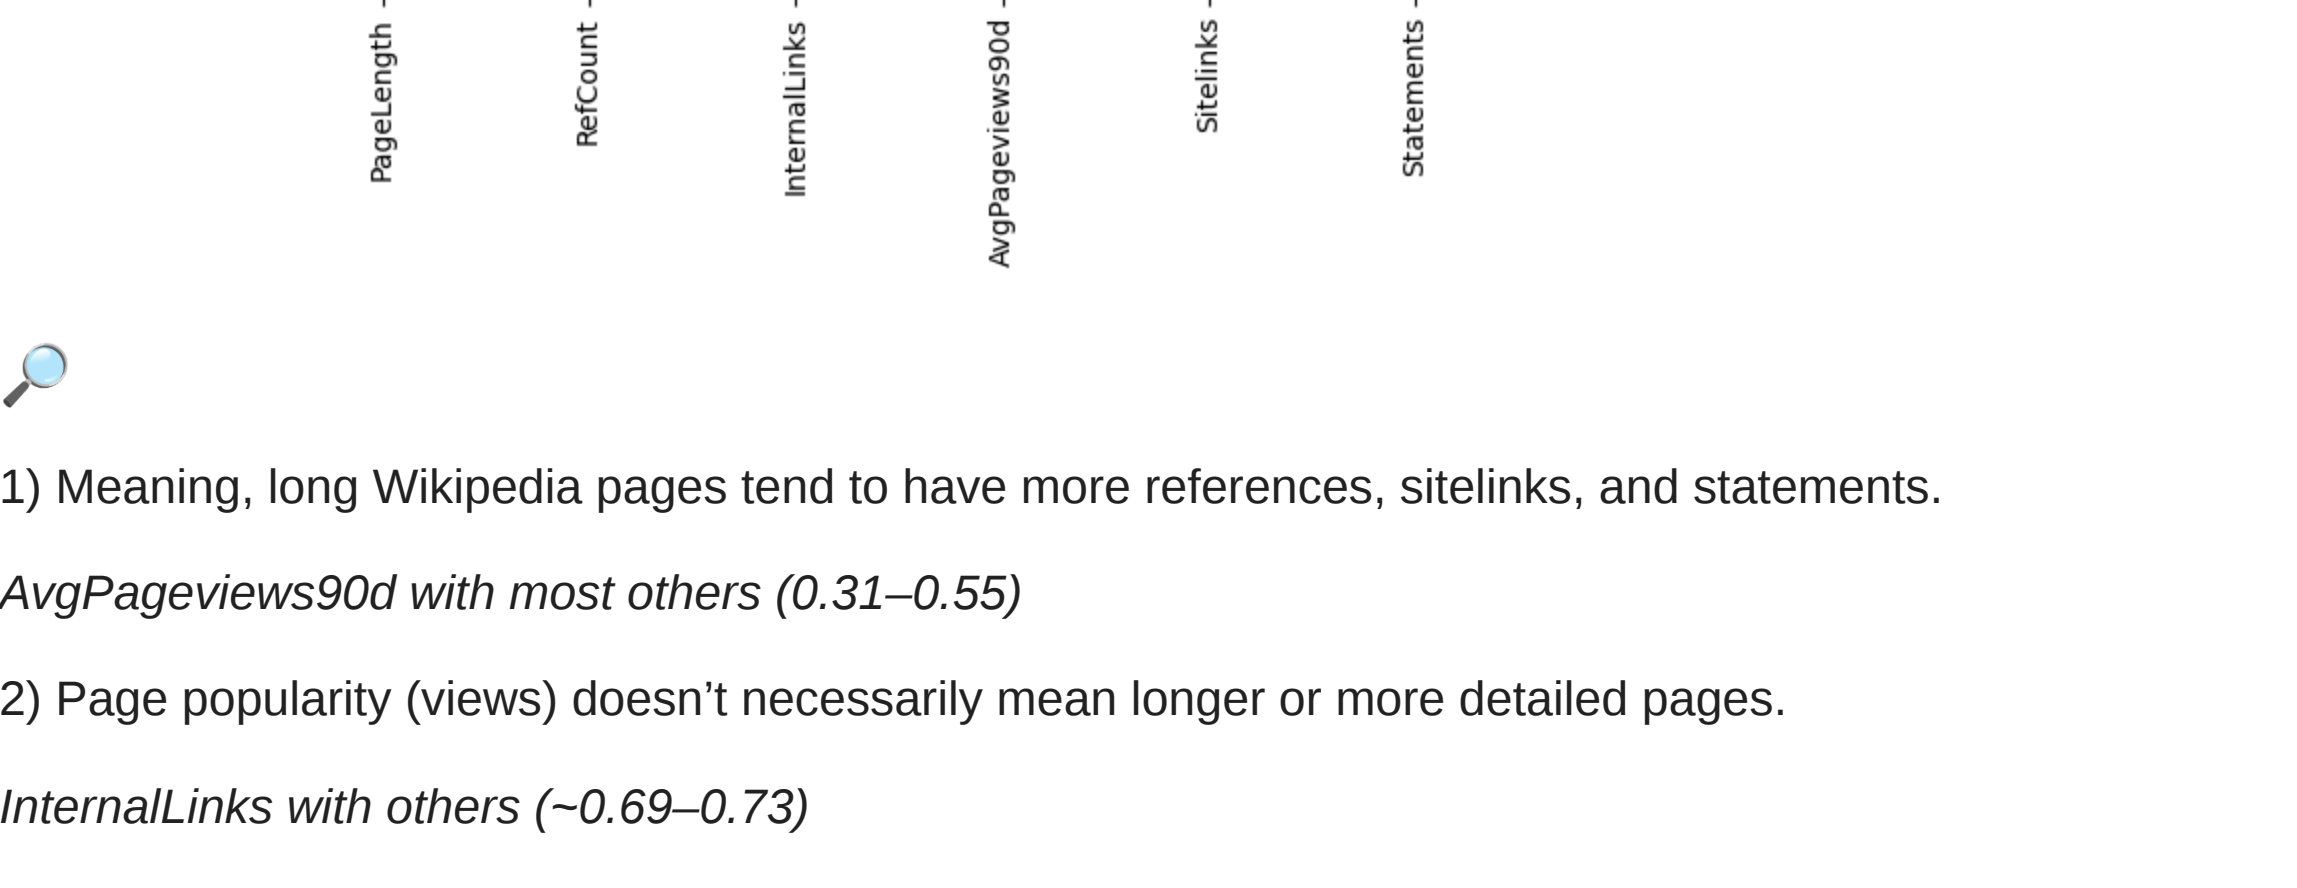

Visualizing at the Brand Level



Wikipedia Page Existence



ChatGPT Ranking Distribution by Wikipedia Page Existence



1) Meaning, long Wikipedia pages tend to have more references, sitelinks, and statements.

AvgPageviews90d with most others (0.31–0.55)

2) Page popularity (views) doesn't necessarily mean longer or more detailed pages.

InternalLinks with others (~0.69–0.73)

3) Pages with more internal links also tend to be longer and better referenced.

Building RichnessScore:

Mathematical Formula for Richness:

log-transform each metric → sum them → get the score.

```
In [8]: def add_logsum_richness(df):
    # defining the Wikipedia-related metrics
    metrics = ["PageLength", "RefCount", "Categories", "InternalLinks", "AvgPageviews90d", "SiteLinks", "Statements"]

    # Ensure numeric conversion (non-numeric -> NaN -> 0)
    for col in metrics:
        df[col] = pd.to_numeric(df[col], errors='coerce')

    # Replacing NaN with 0 (safely) and compute log(1 + value)
    log_vals = np.log1p(df[metrics].fillna(0))

    # Sum across all metrics = RichnessScore
    df["RichnessScore"] = log_vals.sum(axis=1)

    return df

# Apply function
df = add_logsum_richness(df)
print(df[['brand_Name', 'RichnessScore']].head())

brand_Name RichnessScore
0 Samsung 44.576985
1 SanDisk 34.651456
2 Crucial 36.832454
3 Samsung 44.576985
4 LaCie 29.144293
```

Feature engineering

Mention rate per brand across prompts-

Top-K label for ranking analysis-

Average/best rank per brand-

```
In [14]: TOTAL_PROMPTS = df['Prompt'].nunique()

brand_mentions = {
    df.groupby('brand_Name')
    .agg(Mentions=('Prompt', 'nunique'))
    .reset_index()
}

brand_mentions['MentionRate'] = brand_mentions['Mentions'] / TOTAL_PROMPTS

K = 3
df['TopK'] = (df['Rank'] <= K).astype(int)

brand_rank = {
    df.groupby('brand_Name')
    .agg(AvgRank=('Rank', 'mean'),
    TopK_rate=('TopK', 'mean'))
    .reset_index()
}

brand_richness = {
    df.groupby('brand_Name')
    .agg(RichnessScore=('RichnessScore', 'median'),
    PageExists=('PageExists', 'max'))
    .reset_index()
}

brand_level = {
    brand_mentions,
    merge(brand_mentions, on='brand_Name', how='left'),
    merge(brand_rank, on='brand_Name', how='left')
}.fillna({'MentionRate':0, 'AvgRank':np.nan, 'TopK_rate':0})

brand_mentions.sort_values('Mentions', ascending=False)
```

	brand_Name	Mentions	MentionRate
1776	Samsung	10	0.03003
2064	Toyota	9	0.02903
909	Honda	7	0.023102
112	Apple	7	0.023102
1112	LG	7	0.023102
...	...	...	...
783	Galaxy S25 Plus	1	0.003300
782	Galaxy S25	1	0.003300
781	Gainful Grass-Fed Collagen	1	0.003300
780	Gainful	1	0.003300
2287	Edel Chocolate	1	0.003300
2288 rows × 4 columns			

	brand_Name	AvgRank	TopK_rate
0	1688.com	6.0	0.0
1	1MORE	3.0	1.0
2	1-800 Contacts	1.0	1.0
3	2024 Parini Absolute Football Flat Pack Box	15.0	0.0
4	2024 Parini Clearly Dornuss Football Flat Pack Box	14.0	0.0
...	...	...	...
283	iRobot Roomba Combo j+	6.0	0.0
2284	QUIP	1.5	1.0
2285	sous-vide immersion circulators	NaN	0.0
2286	thermometer	NaN	0.0
2287	Edel chocolate	NaN	0.0
2288 rows × 4 columns			

```
In [19]: subset = df[['brand_Name', 'RichnessScore', 'Rank']]
unique_brands = subset.drop_duplicates(subset=['brand_Name']).reset_index(drop=True)
print(unique_brands.head())

brand_Name RichnessScore Rank
0 Samsung 44.576985 1.0
1 SanDisk 34.651456 2.0
2 Crucial 36.832454 3.0
3 LaCie 29.144293 NaN
4 Kingston 34.155925 NaN
```

6) Primary tests

6.1 Correlation: Richness vs MentionRate (brand level)

```
In [27]: # Only keep rows where PageExists is True
subset = brand_level[brand_level['PageExists'] == True]

# Spearman correlation between RichnessScore and MentionRate
rho, p = spearmanr(subset['RichnessScore'], subset['MentionRate'])
print(f"Spearman ρ (Richness vs MentionRate): {rho:.3f}, p={p:.4f}")

Spearman ρ (Richness vs MentionRate): 0.049, p=0.0633
```

👉 This means: Among brands that have a Wikipedia page, richness does not significantly predict how often they are mentioned by ChatGPT.

6.2 Correlation: Richness vs AvgRank (brand level)

```
In [28]: # Subset only brands that actually have AvgRank values
subset = brand_level[brand_level['AvgRank'].notna()]

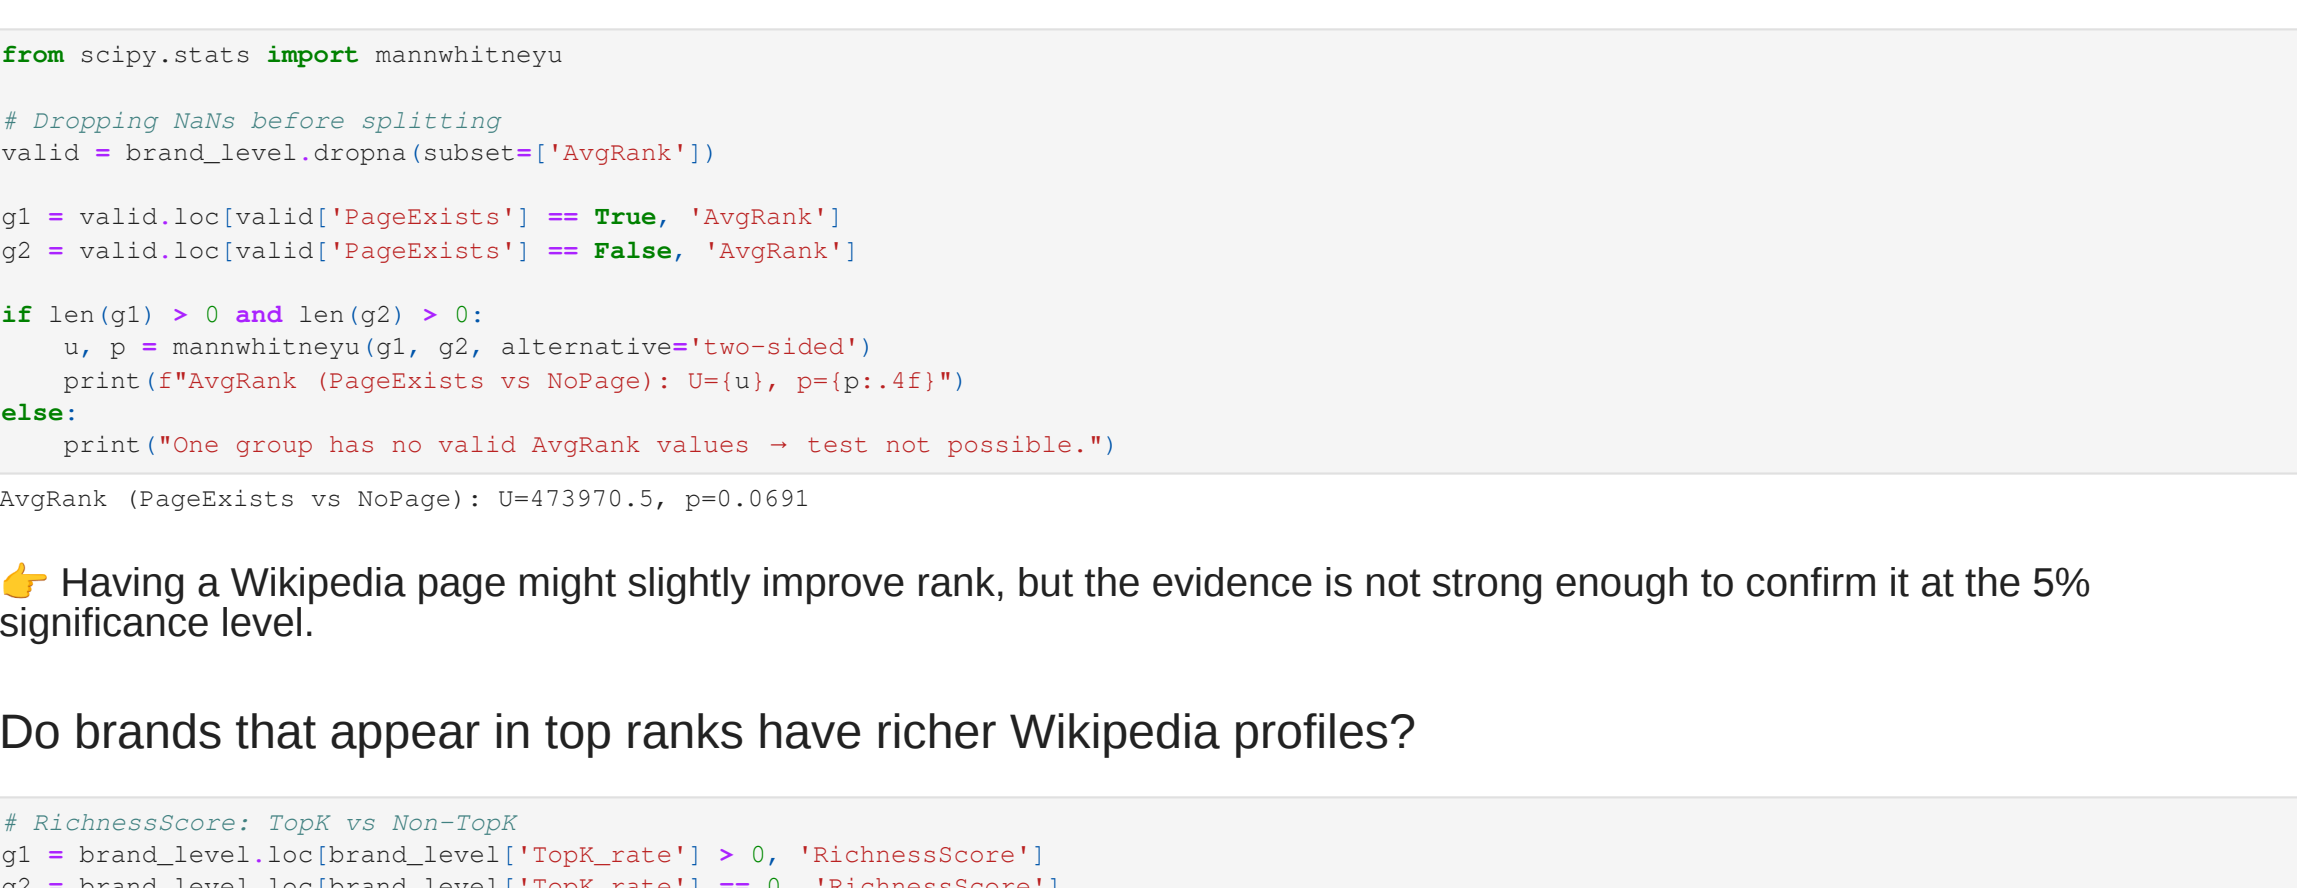
# Spearman correlation between RichnessScore and AvgRank
rho_r, p_r = spearmanr(subset['RichnessScore'], subset['AvgRank'])
print(f"Spearman ρ (Richness vs AvgRank): {rho_r:.3f}, p={p_r:.4f}")

Spearman ρ (Richness vs AvgRank): -0.049, p=0.0266
```

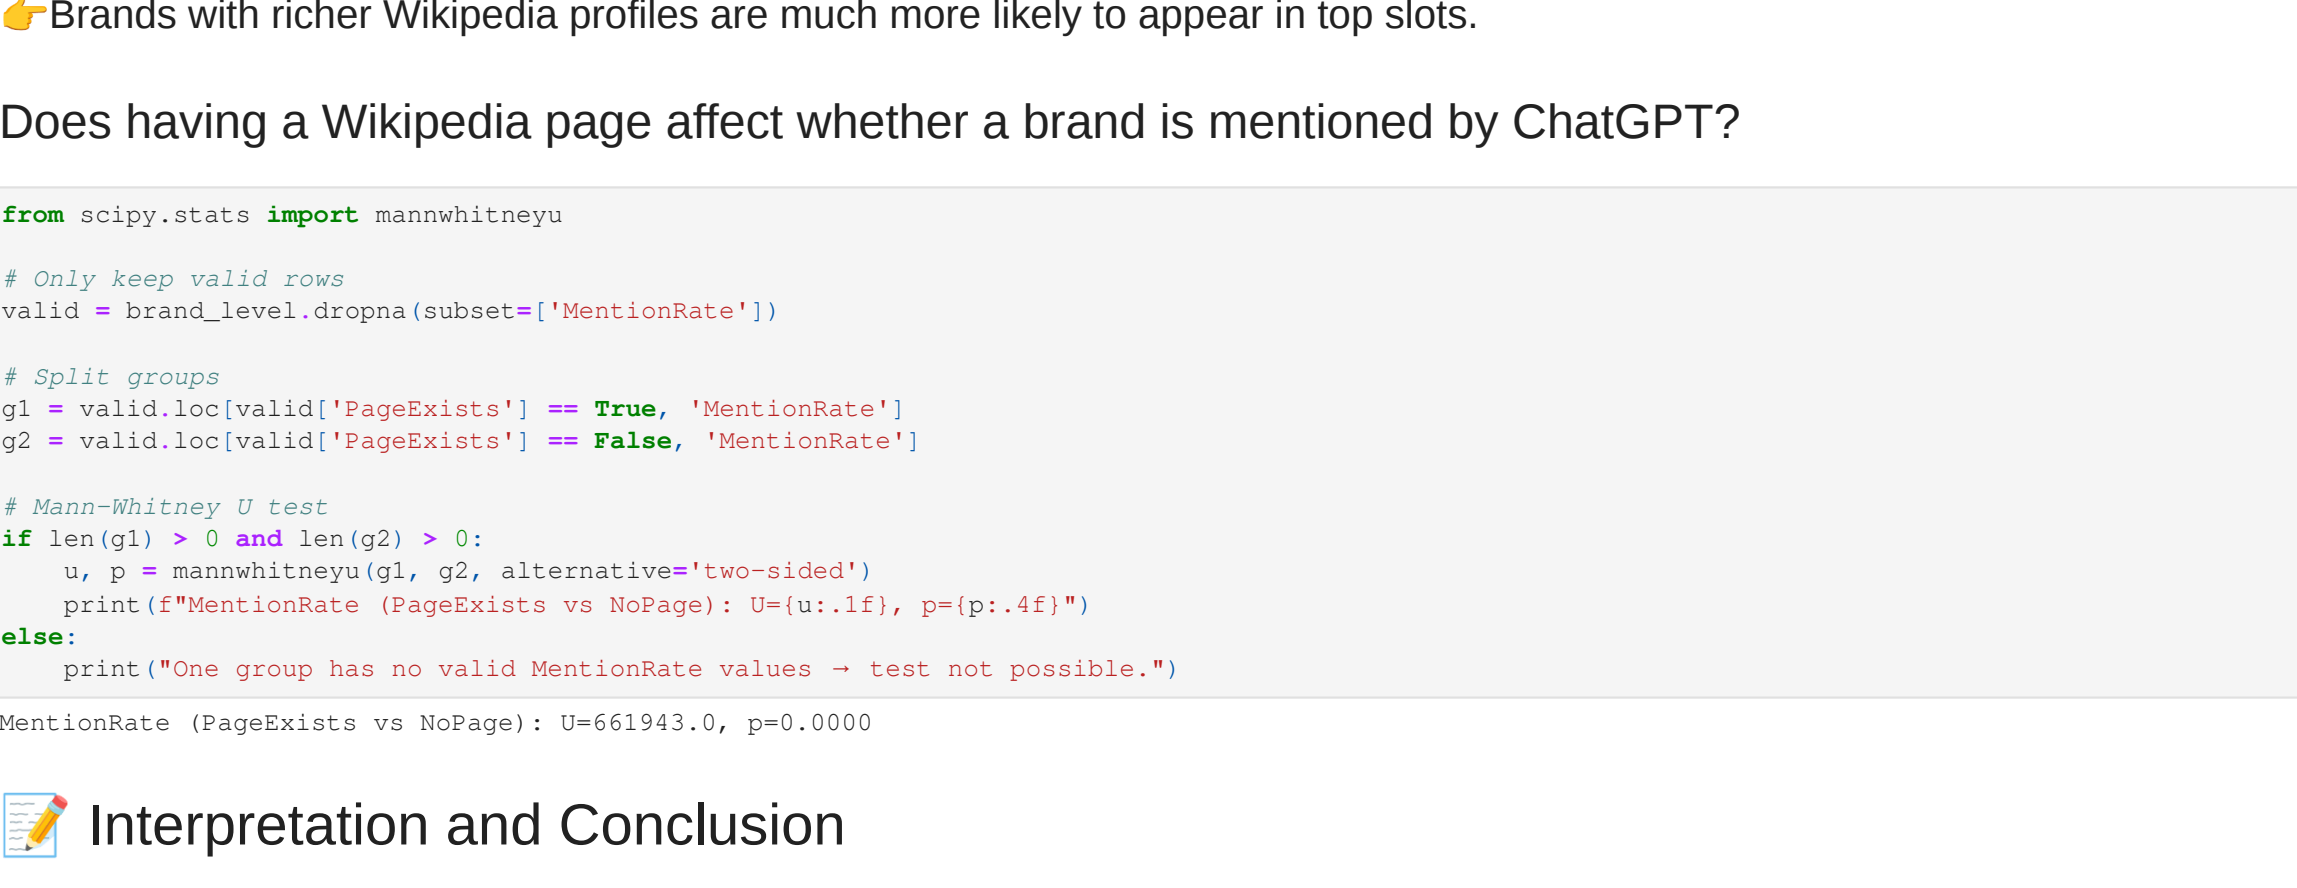
👉 Meaning: Brands with richer Wikipedia/Wikidata profiles tend to have slightly better (lower) average ranks when mentioned by ChatGPT.

But the effect size ($\rho = -0.05$) is so small that it's not practically meaningful.

Richness vs Mention Rate



Richness vs Average Rank in ChatGPT Responses



Do brands with Wikipedia pages have better ranks?

```
In [69]: from scipy.stats import mannwhitneyu

# Dropping NAs before splitting
valid = brand_level.dropna(subset=['AvgRank'])

q1 = valid.loc(valid['PageExists'] == True, 'AvgRank')
q2 = valid.loc(valid['PageExists'] == False, 'AvgRank')

if len(q1) > 0 and len(q2) > 0:
    u, p = mannwhitneyu(q1, q2, alternative='two-sided')
else:
    print("One group has no valid AvgRank values -> test not possible.")

AvgRank (PageExists vs NoPage): U=673970.0, p=0.0691

👉 Having a Wikipedia page might slightly improve rank, but the evidence is not strong enough to confirm it at the 5% significance level.
```

Do brands that appear in top ranks have richer Wikipedia profiles?

```
In [90]: # RichnessScore: TopK vs Non-TopK
q1 = brand_level.loc(brand_level['TopK_rate'] > 0, 'RichnessScore')
q2 = brand_level.loc(brand_level['TopK_rate'] == 0, 'RichnessScore')

u, p = mannwhitneyu(q1, q2, alternative='two-sided')
print(f"Spearman ρ (TopK vs Non-TopK): U={u}, p={p:.4f}")

RichnessScore (TopK vs Non-TopK): U=668909.0, p=0.0000

👉 Brands with richer Wikipedia profiles are much more likely to appear in top slots.
```

Does having a Wikipedia page affect whether a brand is mentioned by ChatGPT?

```
In [96]: from scipy.stats import mannwhitneyu

# Only keep valid rows
valid = brand_level.dropna(subset=['MentionRate'])

# Split groups
q1 = valid.loc(valid['PageExists'] == True, 'MentionRate')
q2 = valid.loc(valid['PageExists'] == False, 'MentionRate')

# Mann-Whitney U test
if len(q1) > 0 and len(q2) > 0:
    u, p = mannwhitneyu(q1, q2, alternative='two-sided')
else:
    print("One group has no valid MentionRate values -> test not possible.")

MentionRate (PageExists vs NoPage): U=661943.0, p=0.0000
```

Interpretation and Conclusion

1. Richness vs. Top-K Presence

3Test: Mann–Whitney U

Result: U = 668909.0, $p < 0.0001$

Interpretation: Brands with richer Wikipedia/Wikidata profiles are much more likely to appear in the top-ranked slots of ChatGPT responses.

Implication: Having a strong Wikipedia/Wikidata presence can boost the chances of being highlighted prominently.

2. Average Rank (PageExists vs. NoPage)

Test: Mann–Whitney U

Result: U = 473970.5, $p = 0.0691$

Interpretation: Brands with a Wikipedia page tend to rank slightly better on average, but the evidence is not statistically significant at the 5% threshold.

Implication: Merely having a Wikipedia page may help visibility, but the effect is weak and uncertain.

3. Richness vs. Average Rank

Test: Spearman correlation

Result: $\rho = -0.049$, $p = 0.0266$

Interpretation: There is a very weak negative correlation, meaning richer profiles are associated with slightly better (lower) average ranks.

Implication: Although statistically significant, the effect size is so small that it has little practical importance.

4. Richness vs. Mention Rate

Test: Spearman correlation

Result: $\rho = 0.049$, $p = 0.0633$

Interpretation: Among brands with Wikipedia pages, richness is not a significant predictor of how often they are mentioned by ChatGPT.

Implication: Richer profiles don't guarantee higher frequency of mentions.

Overall Conclusion

Strong evidence:

1) Brands with richer profiles are more likely to appear in top-ranked slots.

2) Having a Wikipedia page strongly increases the chances of being mentioned by ChatGPT across prompts. This confirms that Wikipedia presence itself (even without richness details) matters a lot for visibility.

Minimal evidence: Richness has almost no practical effect on how often brands are mentioned, nor on their exact rank when mentioned.

👉 Takeaway: Investing in richer Wikipedia/Wikidata profiles improves the chance of being placed in top slots (visibility impact). However, it does not guarantee more frequent mentions or meaningfully better rankings once a brand is included.

```
In [ ]:
```

```
In [ ]:
```